

# Predicting the Income Category (Census-Based) Using Different Classifiers - Parth Shroff

August 16, 2020

## 1 Introduction to the dataset

The dataset comes from the UC Irvine database. The extraction was done by Barry Becker from the 1994 Census Database. The prediction task is the income category (greater than \$50K or less than \$50K). Features include Age, Class of Work, Final Weight (Estimate of the number of people the census believes the entry is generalized to), Education, Education by Year, Marital Status, Occupation, Relationship, Race, Sex, Capital-Gain, Capital-Loss, Hours per week, and Native Country.

## 2 Implementation of Classifiers

Import statements from pandas, seaborn, matplotlib, and sklearn.

```
[37]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC #Support Vector Classifier
from sklearn import svm #Support Vector Machine
from sklearn.neural_network import MLPClassifier
from sklearn.metrics import confusion_matrix, classification_report
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split, cross_val_score, GridSearchCV, RandomizedSearchCV
from sklearn.metrics import accuracy_score, f1_score, make_scorer
from sklearn.tree import DecisionTreeClassifier
%matplotlib inline
```

### 2.1 Reformating and Analyzing Data

```
[2]: income = pd.read_csv('adult.csv')
```

```
[3]: income.head()
```

```
[3]:
```

	Age	Class of Work	Final Weight	Education	Education by Year	\
0	39	State-gov	77516	Bachelors	13	
1	50	Self-emp-not-inc	83311	Bachelors	13	
2	38	Private	215646	HS-grad	9	
3	53	Private	234721	11th	7	
4	28	Private	338409	Bachelors	13	

	Marital Status	Occupation	Relationship	Race	Sex	\
0	Never-married	Adm-clerical	Not-in-family	White	Male	
1	Married-civ-spouse	Exec-managerial	Husband	White	Male	
2	Divorced	Handlers-cleaners	Not-in-family	White	Male	
3	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	
4	Married-civ-spouse	Prof-specialty	Wife	Black	Female	

	Capital-Gain	Capital-Loss	Hours per week	Native Country	Income Category
0	2174	0	40	United-States	<=50K
1	0	0	13	United-States	<=50K
2	0	0	40	United-States	<=50K
3	0	0	40	United-States	<=50K
4	0	0	40	Cuba	<=50K

```
[4]: income.describe()
```

```
[4]:
```

	Age	Final Weight	Education by Year	Capital-Gain	\
count	32561.000000	3.256100e+04	32561.000000	32561.000000	
mean	38.581647	1.897784e+05	10.080679	1077.648844	
std	13.640433	1.055500e+05	2.572720	7385.292085	
min	17.000000	1.228500e+04	1.000000	0.000000	
25%	28.000000	1.178270e+05	9.000000	0.000000	
50%	37.000000	1.783560e+05	10.000000	0.000000	
75%	48.000000	2.370510e+05	12.000000	0.000000	
max	90.000000	1.484705e+06	16.000000	99999.000000	

	Capital-Loss	Hours per week
count	32561.000000	32561.000000
mean	87.303830	40.437456
std	402.960219	12.347429
min	0.000000	1.000000
25%	0.000000	40.000000
50%	0.000000	40.000000
75%	0.000000	45.000000
max	4356.000000	99.000000

```
[5]: income.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
```

Data columns (total 15 columns):

#	Column	Non-Null Count	Dtype
0	Age	32561 non-null	int64
1	Class of Work	32561 non-null	object
2	Final Weight	32561 non-null	int64
3	Education	32561 non-null	object
4	Education by Year	32561 non-null	int64
5	Marital Status	32561 non-null	object
6	Occupation	32561 non-null	object
7	Relationship	32561 non-null	object
8	Race	32561 non-null	object
9	Sex	32561 non-null	object
10	Capital-Gain	32561 non-null	int64
11	Capital-Loss	32561 non-null	int64
12	Hours per week	32561 non-null	int64
13	Native Country	32561 non-null	object
14	Income Category	32561 non-null	object

dtypes: int64(6), object(9)

memory usage: 3.7+ MB

```
[6]: income['Income Category'].replace([' <=50K', ' >50K'], [0, 1], inplace = True)
income['Class of Work'].replace([' Private', ' Self-emp-not-inc', '
↳Self-emp-inc', ' Federal-gov', ' Local-gov', ' State-gov', ' Without-pay', '
↳Never-worked'], [0, 1, 2, 3, 4, 5, 6, 7], inplace = True)
income['Marital Status'].replace(['Married-civ-spouse', 'Divorced',
↳'Never-married', 'Separated', 'Widowed', 'Married-spouse-absent',
↳'Married-AF-spouse'], [0, 1, 2, 3, 4, 5, 6], inplace = True)
income['Sex'].replace(['Male', 'Female'], [0, 1], inplace = True)
income['Occupation'].replace(['Tech-support', 'Craft-repair', 'Other-service',
↳'Sales', 'Exec-managerial', 'Prof-specialty', 'Handlers-cleaners',
↳'Machine-op-inspct', 'Adm-clerical', 'Farming-fishing', 'Transport-moving',
↳'Priv-house-serv', 'Protective-serv', 'Armed-Forces'], [0, 1, 2, 3, 4, 5, 6,
↳7, 8, 9, 10, 11, 12, 13], inplace = True)
income['Education'].replace(['Bachelors', 'Some-college', '11th', 'HS-grad',
↳'Prof-school', 'Assoc-acdm', 'Assoc-voc', '9th', '7th-8th', '12th',
↳'Masters', '1st-4th', '10th', 'Doctorate', '5th-6th', 'Preschool'], [0, 1,
↳2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15], inplace = True)
income['Relationship'].replace(['Wife', 'Own-child', 'Husband',
↳'Not-in-family', 'Other-relative', 'Unmarried'], [0, 1, 2, 3, 4, 5], inplace
↳= True)
income['Race'].replace(['White', 'Asian-Pac-Islander', 'Amer-Indian-Eskimo',
↳'Other', 'Black'], [0, 1, 2, 3, 4], inplace = True)
```

```
income['Native Country'].replace(['United-States', 'Cambodia', 'England',
↳ 'Puerto-Rico', 'Canada', 'Germany', 'Outlying-US(Guam-USVI-etc)', 'India',
↳ 'Japan', 'Greece', 'South', 'China', 'Cuba', 'Iran', 'Honduras',
↳ 'Philippines', 'Italy', 'Poland', 'Jamaica', 'Vietnam', 'Mexico',
↳ 'Portugal', 'Ireland', 'France', 'Dominican-Republic', 'Laos', 'Ecuador',
↳ 'Taiwan', 'Haiti', 'Columbia', 'Hungary', 'Guatemala', 'Nicaragua',
↳ 'Scotland', 'Thailand', 'Yugoslavia', 'El-Salvador', 'Trinidad&Tobago',
↳ 'Peru', 'Hong', 'Holand-Netherlands'], [0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10,
↳ 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29,
↳ 30, 31, 32, 33, 34, 35, 36, 37, 38, 39, 40], inplace = True)
```

```
[7]: income = income[income['Class of Work'] != '?']
income = income[income['Occupation'] != '?']
income = income[income['Native Country'] != '?']
```

```
[8]: income.head(10)
```

```
[8]:   Age Class of Work  Final Weight  Education  Education by Year  \
0    39              5         77516           0              13
1    50              1         83311           0              13
2    38              0        215646           3              9
3    53              0        234721           2              7
4    28              0        338409           0             13
5    37              0        284582          10             14
6    49              0        160187           7              5
7    52              1        209642           3              9
8    31              0         45781          10             14
9    42              0        159449           0             13
```

```
   Marital Status Occupation  Relationship  Race  Sex  Capital-Gain  \
0              2           8              3    0    0          2174
1              0           4              2    0    0              0
2              1           6              3    0    0              0
3              0           6              2    4    0              0
4              0           5              0    4    1              0
5              0           4              0    0    1              0
6              5           2              3    4    1              0
7              0           4              2    0    0              0
8              2           5              3    0    1         14084
9              0           4              2    0    0          5178
```

```
   Capital-Loss  Hours per week Native Country  Income Category
0              0              40              0              0
1              0              13              0              0
2              0              40              0              0
3              0              40              0              0
4              0              40             12              0
```

5	0	40	0	0
6	0	16	18	0
7	0	45	0	1
8	0	50	0	1
9	0	40	0	1

## 2.2 Data Visualization

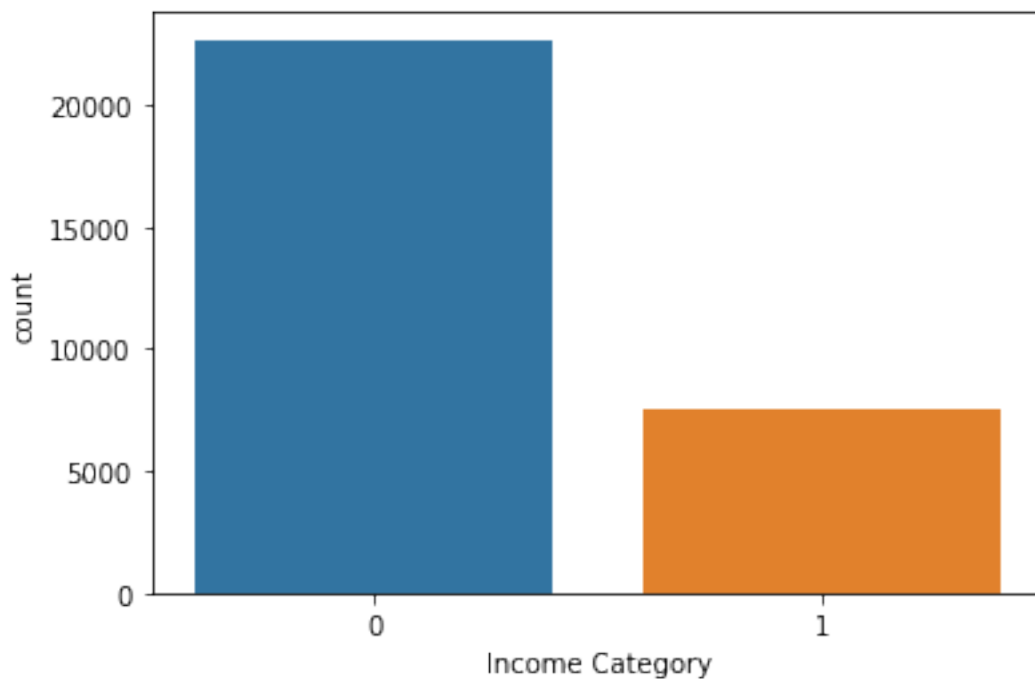
I chose to visualize some of the features and the respective frequency of the income category to better understand the dataset.

```
[9]: income['Income Category'].value_counts()
```

```
[9]: 0    22654
     1     7508
     Name: Income Category, dtype: int64
```

```
[10]: sns.countplot(income['Income Category'])
```


```
[10]: <matplotlib.axes._subplots.AxesSubplot at 0x1b246672788>
```



```
[11]: plt.hist(income[income["Income Category"]==1].Age.values, 10, facecolor='blue',
             ↪alpha=0.5)
       plt.title("Age distribution with above $50K annual income")
       plt.xlabel("Age")
```

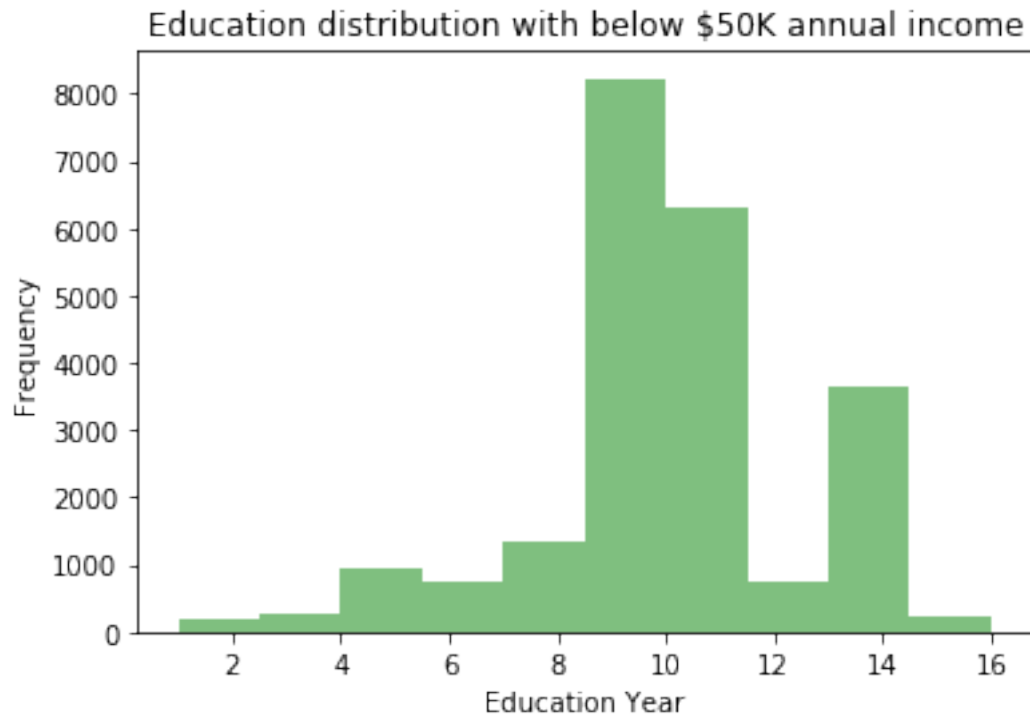
```
plt.ylabel("Frequency")
```

```
[11]: Text(0, 0.5, 'Frequency')
```



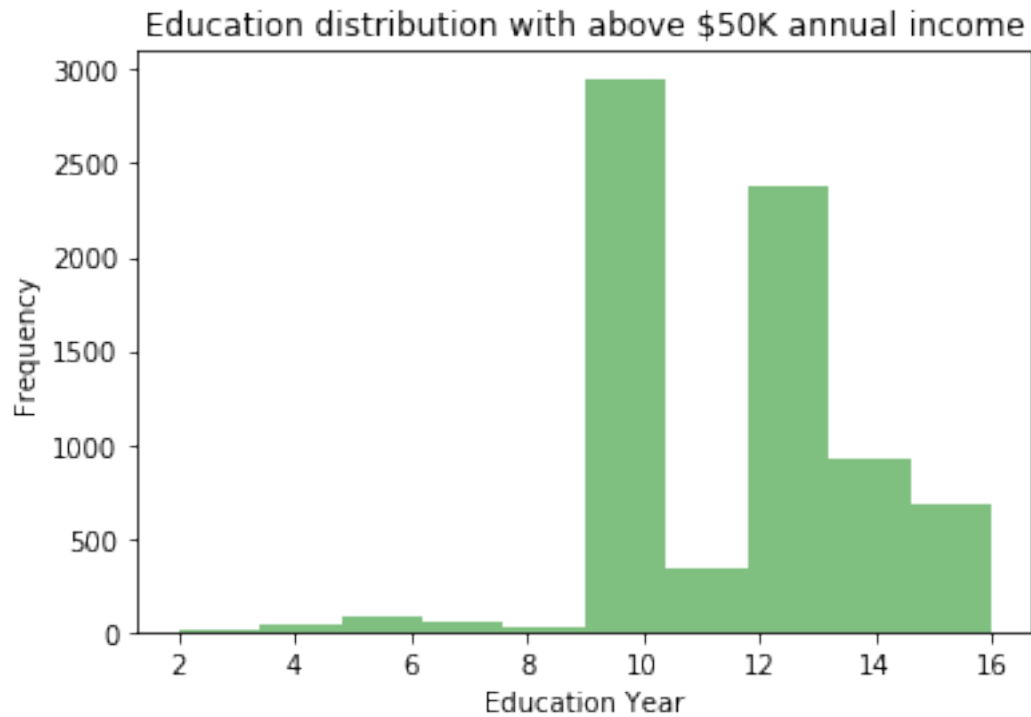
```
[59]: plt.hist(income[income["Income Category"]==0]["Education by Year"].values, 10,
    ↪facecolor='green', alpha=0.5)
plt.title("Education distribution with below $50K annual income")
plt.xlabel("Education Year")
plt.ylabel("Frequency")
```

```
[59]: Text(0, 0.5, 'Frequency')
```



```
[61]: plt.hist(income[income["Income Category"]==1]["Education by Year"].values, 10,
             ↳facecolor='green', alpha=0.5)
plt.title("Education distribution with above $50K annual income")
plt.xlabel("Education Year")
plt.ylabel("Frequency")
```


```
[61]: Text(0, 0.5, 'Frequency')
```



```
[62]: plt.hist(income[income["Income Category"]==0]["Hours per week"].values, 10,
             ↳facecolor='red', alpha=0.5)
plt.title("Hours/week distribution with below $50K annual income")
plt.xlabel("Hours/week")
plt.ylabel("Frequency")
```


```
[62]: Text(0, 0.5, 'Frequency')
```





```
[63]: plt.hist(income[income["Income Category"]==1]["Hours per week"].values, 10,
             ↳facecolor='red', alpha=0.5)
plt.title("Hours/week distribution with above $50K annual income")
plt.xlabel("Hours/week")
plt.ylabel("Frequency")
```

```
[63]: Text(0, 0.5, 'Frequency')
```



### 2.3 Splitting the Dataset into a Training and Testing Set.

The `train_test_split` method from the `model_selection` package in `sklearn` ensures that both the features and the output (income category) are split. I chose to have the test size be 35% of the dataset allowing for 65% of the dataset to train the classifier

```
[13]: X = income.drop('Income Category', axis=1) #All of the features besides quality
      y = income['Income Category']
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.25,
      ↪random_state = 10)
```

I applied a Standard Scaler to the data to ensure that the data is normally distributed and  $\mu = 0$  and  $\sigma = 1$ . This is important because features such as *Final Weight* and *Capital Gain* have much larger values than *Relationship*, or *Race*. The model should not be skewed in favor of the large parameters. Each feature should be treated equally.

```
[14]: sc = StandardScaler()
      X_train = sc.fit_transform(X_train)
      X_test = sc.transform(X_test)
      X_train[:5]
```

```
[14]: array([[ -1.63666649,  -0.49894997,  -0.22409534,  -0.40161329,  -1.21756042,
           0.81245913,   0.42499252,  -1.13607777,  -0.37238476,  -0.69180443,
          -0.14686997,  -0.21743312,  -1.73886654,   6.46837417],
```

```

[-0.33997058, -0.49894997,  0.75448338, -0.6932351 , -0.04239149,
 0.81245913,  1.09592251, -1.13607777, -0.37238476,  1.44549522,
-0.14686997, -0.21743312, -0.07490793, -0.26843863],
[ 0.04141056,  0.18800969,  1.39424134,  1.9313612 ,  1.52450043,
-0.89726113, -1.25233245, -0.32201286, -0.37238476, -0.69180443,
-0.14686997, -0.21743312,  1.17306103, -0.26843863],
[-0.33997058, -0.49894997,  0.18818156, -0.6932351 , -0.04239149,
-0.89726113, -1.58779745, -0.32201286,  2.9591223 , -0.69180443,
-0.14686997, -0.21743312, -0.07490793, -0.26843863],
[ 0.42279171,  2.24888869, -0.49810685,  1.9313612 ,  1.52450043,
-0.89726113,  0.08952753, -0.32201286, -0.37238476, -0.69180443,
 1.94914465, -0.21743312, -0.49089758, -0.26843863]])

```

## 2.4 Random Forest Classifier

```

[15]: RFC = RandomForestClassifier(n_estimators=150)
      RFC.fit(X_train, y_train)
      pred_RFC = RFC.predict(X_test)

```

```

[16]: print(classification_report(y_test, pred_RFC))

```

	precision	recall	f1-score	support
0	0.88	0.93	0.90	5656
1	0.75	0.62	0.68	1885
accuracy			0.85	7541
macro avg	0.81	0.78	0.79	7541
weighted avg	0.85	0.85	0.85	7541

```

[18]: print(confusion_matrix(y_test, pred_RFC))

```

```

[[5264  392]
 [ 716 1169]]

```

```

[20]: acc = accuracy_score(y_test, pred_RFC)
      acc

```

```

[20]: 0.8530698846306856

```

## 2.5 Support Vector Machine Classifier

```

[21]: clf = svm.SVC()
      clf.fit(X_train, y_train)
      pred_clf = clf.predict(X_test)

```

```
[22]: print(classification_report(y_test, pred_clf))
```

	precision	recall	f1-score	support
0	0.86	0.94	0.90	5656
1	0.76	0.56	0.64	1885
accuracy			0.85	7541
macro avg	0.81	0.75	0.77	7541
weighted avg	0.84	0.85	0.84	7541

```
[23]: print(confusion_matrix(y_test, pred_clf))
```

```
[[5324  332]
 [ 833 1052]]
```

```
[53]: acc = accuracy_score(y_test, pred_clf)
acc
```

```
[53]: 0.845511205410423
```

## 2.6 Multi-layer Perceptron Classifier

```
[25]: MLP = MLPClassifier(hidden_layer_sizes = (14, 14, 14), max_iter=1000)
MLP.fit(X_train, y_train)
pred_MLP = MLP.predict(X_test)
```

```
[26]: print(classification_report(y_test, pred_MLP))
```

	precision	recall	f1-score	support
0	0.89	0.91	0.90	5656
1	0.71	0.66	0.68	1885
accuracy			0.85	7541
macro avg	0.80	0.78	0.79	7541
weighted avg	0.84	0.85	0.84	7541

```
[27]: print(confusion_matrix(y_test, pred_MLP))
```

```
[[5153  503]
 [ 650 1235]]
```

```
[28]: acc = accuracy_score(y_test, pred_MLP)
acc
```

```
[28]: 0.8471025062988994
```

## 2.7 Decision Tree Classifier

```
[29]: DTC = DecisionTreeClassifier(random_state=100, max_depth=5)
      DTC.fit(X_train, y_train)
```

```
[29]: DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None, criterion='gini',
                             max_depth=5, max_features=None, max_leaf_nodes=None,
                             min_impurity_decrease=0.0, min_impurity_split=None,
                             min_samples_leaf=1, min_samples_split=2,
                             min_weight_fraction_leaf=0.0, presort='deprecated',
                             random_state=100, splitter='best')
```

```
[30]: train_predictions = DTC.predict(X_train)
      pred_DTC = DTC.predict(X_test)
      print(confusion_matrix(y_test, pred_DTC))
```

```
[[5385  271]
 [ 898  987]]
```

```
[31]: print(classification_report(y_test, pred_DTC))
```

	precision	recall	f1-score	support
0	0.86	0.95	0.90	5656
1	0.78	0.52	0.63	1885
accuracy			0.84	7541
macro avg	0.82	0.74	0.77	7541
weighted avg	0.84	0.84	0.83	7541

```
[32]: acc = accuracy_score(y_test, pred_DTC)
      acc
```

```
[32]: 0.8449807717809309
```

## 2.8 Hyperparameter Tuning (Support Vector Machine)

```
[33]: score = cross_val_score(tree_clf, X_train, y_train, cv = 5, scoring =
      ↪ 'f1_macro')
      score.mean()
```

```
[33]: 0.7559805973263133
```

```
[34]: #Checking to see if the Decision Tree is overfitted.
print("The Training F1 Score is", f1_score(train_predictions, y_train))
print("The Testing F1 Score is", f1_score(pred_dtc, y_test))
```

The Training F1 Score is 0.6188070929607739

The Testing F1 Score is 0.6280623608017818

The results below for the GridSearchCV demonstrate the scores from the five different cross-validation trials. The RBF kernel with a C value of 1 seemed to produce the best mean test score but each of the five scores is still very close to each other.

```
[47]: GSCV = GridSearchCV(svm.SVC(gamma = "auto"), {
    'C': [1, 10, 20],
    'kernel': ['rbf', 'linear']
}, cv = 5, return_train_score = False)

GSCV.fit(X_train, y_train)
```

```
[47]: GridSearchCV(cv=5, error_score=nan,
    estimator=SVC(C=1.0, break_ties=False, cache_size=200,
    class_weight=None, coef0=0.0,
    decision_function_shape='ovr', degree=3,
    gamma='auto', kernel='rbf', max_iter=-1,
    probability=False, random_state=None, shrinking=True,
    tol=0.001, verbose=False),
    iid='deprecated', n_jobs=None,
    param_grid={'C': [1, 10, 20], 'kernel': ['rbf', 'linear']},
    pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
    scoring=None, verbose=0)
```

```
[49]: GSCV.cv_results_
```

```
[49]: {'mean_fit_time': array([ 9.44090753, 11.50188084, 13.94156656, 40.65536585,
    16.23147578,
    71.63549361]),
    'std_fit_time': array([0.81035446, 0.13484138, 0.21015353, 0.80159731,
    1.27727564,
    4.04634212]),
    'mean_score_time': array([1.33945813, 0.69874735, 1.36595359, 0.69614177,
    1.3671536 ,
    0.70673208]),
    'std_score_time': array([0.12649319, 0.00381951, 0.00560563, 0.0031446 ,
    0.0054026 ,
    0.01420667]),
    'param_C': masked_array(data=[1, 1, 10, 10, 20, 20],
    mask=[False, False, False, False, False, False],
    fill_value='?',
    dtype=object),
```

```

'param_kernel': masked_array(data=['rbf', 'linear', 'rbf', 'linear', 'rbf',
'linear'],
                             mask=[False, False, False, False, False, False],
                             fill_value='?',
                             dtype=object),
'params': [{ 'C': 1, 'kernel': 'rbf'},
            { 'C': 1, 'kernel': 'linear'},
            { 'C': 10, 'kernel': 'rbf'},
            { 'C': 10, 'kernel': 'linear'},
            { 'C': 20, 'kernel': 'rbf'},
            { 'C': 20, 'kernel': 'linear'}],
'split0_test_score': array([0.84287293, 0.83911602, 0.84508287, 0.83911602,
0.84618785,
0.83911602]),
'split1_test_score': array([0.83819629, 0.83443855, 0.83841733, 0.83399646,
0.83952255,
0.83399646]),
'split2_test_score': array([0.84703802, 0.84350133, 0.8428382 , 0.84350133,
0.84173298,
0.84350133]),
'split3_test_score': array([0.84504863, 0.84106985, 0.84328028, 0.84106985,
0.84040672,
0.84106985]),
'split4_test_score': array([0.84416446, 0.8362069 , 0.8428382 , 0.83642794,
0.84173298,
0.83642794]),
'mean_test_score': array([0.84346406, 0.83886653, 0.84249138, 0.83882232,
0.84191661,
0.83882232]),
'std_test_score': array([0.00296221, 0.00325819, 0.00219894, 0.00334821,
0.00229444,
0.00334821]),
'rank_test_score': array([1, 4, 2, 5, 3, 5])}

```

The results below for the RandomSearchCV demonstrate the scores from the two iterations. The RBF kernel with a C value of 1 seemed to produce the best mean test score. I chose to use two iterations to speed up the process.

```

[44]: RS = RandomizedSearchCV(svm.SVC(gamma = "auto"), {
    'C': [1, 10, 20],
    'kernel': ['rbf', 'linear']
}, cv = 5, return_train_score = False, n_iter=2)

RS.fit(X_train, y_train)

```

```

[44]: RandomizedSearchCV(cv=5, error_score=nan,
    estimator=SVC(C=1.0, break_ties=False, cache_size=200,

```

```

class_weight=None, coef0=0.0,
decision_function_shape='ovr', degree=3,
gamma='auto', kernel='rbf', max_iter=-1,
probability=False, random_state=None,
shrinking=True, tol=0.001, verbose=False),
iid='deprecated', n_iter=2, n_jobs=None,
param_distributions={'C': [1, 10, 20],
                    'kernel': ['rbf', 'linear']},
pre_dispatch='2*n_jobs', random_state=None, refit=True,
return_train_score=False, scoring=None, verbose=0)

```

```
[45]: RS.cv_results_
```

```

[45]: {'mean_fit_time': array([ 9.87321062, 70.93206229]),
      'std_fit_time': array([0.40554098, 3.75435546]),
      'mean_score_time': array([1.40964532, 0.69953771]),
      'std_score_time': array([0.01248228, 0.00311803]),
      'param_kernel': masked_array(data=['rbf', 'linear'],
                                   mask=[False, False],
                                   fill_value='?',
                                   dtype=object),
      'param_C': masked_array(data=[1, 20],
                              mask=[False, False],
                              fill_value='?',
                              dtype=object),
      'params': [{'kernel': 'rbf', 'C': 1}, {'kernel': 'linear', 'C': 20}],
      'split0_test_score': array([0.84287293, 0.83911602]),
      'split1_test_score': array([0.83819629, 0.83399646]),
      'split2_test_score': array([0.84703802, 0.84350133]),
      'split3_test_score': array([0.84504863, 0.84106985]),
      'split4_test_score': array([0.84416446, 0.83642794]),
      'mean_test_score': array([0.84346406, 0.83882232]),
      'std_test_score': array([0.00296221, 0.00334821]),
      'rank_test_score': array([1, 2])}

```

### 3 Summary and Analysis

For this project, I chose to use the census income dataset from the UC Irvine Dataset Repository. One of the main problems with the dataset was that about 1% of the data was incomplete. Most of my time was spent in cleaning and preparing the dataset for the four different classification models. I chose to run the following four classification models: Random Forest Classifier, Support Vector Machine Classifier, Multi-layer Perceptron (Neural-Network Model), and the Decision Tree Classifier. The Random Forest Classifier produced the highest accuracy (0.8531) while the Decision Tree Classifier produced the lowest accuracy (0.8449). However, all four of the accuracy values were very close to each other (IQR = 0.00649). This is perhaps due to the large size of the data set and the overfitting of models to the data (refer to the low F1 testing score as a potential indicator of overfitting).



As a result, I tried to optimize the classifier using cross-validation techniques present in the GridSearchCV object from the Sklearn Model Selection package. Since the dataset is very large, and GridSearchCV studies every permutation of the parameter list using K-Fold Cross-Validation, I limited the ‘C’ (cross-validation generator) and ‘Kernel’ (Function type) parameters. In both instances, I received the highest mean test score when  $C = 1$ , and the Kernel was set to the RBF function.